

✓ CPSC5610 Assignment 3

CNN Image classifier for the SVHN dataset

Overview

In this assignment, you will create a convolutional neural network that classifies real-world images digits. You will use concepts from what we learned this course in building, training, testing, validating and saving your Tensorflow classifier model.

This is an open-ended project.

What to submit

- Please do this assignment in a folder.
- When you have completed the notebook, you will save this notebook including all the intermediate results/plots.
- Please also print out a pdf file from your finished notebook, and put it in the same folder.
- Before you submit, please move your data folder outside. It is too big to submit.
- Please include your saved models.
- Then zip your whole folder and submit the single zip file.

Start

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
import os
import zipfile
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import scipy.io
```

In this assignment, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. "Reading Digits in Natural Images with Unsupervised Feature Learning". NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

Please only use the pre-cropped datasets in Format 2.

Both `train` and `test` are dictionaries with keys `x` and `y` for the input images and labels respectively.

(20pts) Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

✓ 1) Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.

```
# Mount google drive
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
# Unzip SVHN data
zip_path = "/content/drive/My Drive/Seattle University/CPSC5610 - Artificial Intelligence/Data/svhn_data.zip"
extract_root = "/content"
```

```
with zipfile.ZipFile(zip_path, 'r') as zip_ref:
    zip_ref.extractall(extract_root)
```

```
# Load .mat files
data_directory = os.path.join(extract_root, "svhn_data")
train = scipy.io.loadmat(os.path.join(data_directory, "train_32x32.mat"))
test = scipy.io.loadmat(os.path.join(data_directory, "test_32x32.mat"))
extra = scipy.io.loadmat(os.path.join(data_directory, "extra_32x32.mat"))
```

```
# Subset extra dataset
#subset_size = 200000
#indices = np.random.choice(extra['X'].shape[-1], subset_size, replace=False)
#extra_X_subset = extra['X'][:, :, :, indices]
#extra_y_subset = extra['y'][indices]
```

```
# Transpose images from (32, 32, 3, N) to (N, 32, 32, 3)
train_images = train['X'].transpose(3, 0, 1, 2) # from (32, 32, 3, 73257) to (73257, 32, 32, 3)
test_images = test['X'].transpose(3, 0, 1, 2) # from (32, 32, 3, 26032) to (26032, 32, 32, 3)
extra_images = extra['X'].transpose(3, 0, 1, 2) # from (32, 32, 3, 531131) to (531131, 32, 32, 3)
```

```
# Flatten labels
train_labels = train['y'].flatten() # (73257,)
test_labels = test['y'].flatten() # (26032,)
extra_labels = extra['y'].flatten() # (531131,)
```

train_images.shape # samples, height, width, channels

(73257, 32, 32, 3)

train_labels.shape # samples

(73257,)

test_images.shape # samples, height, width, channels

(26032, 32, 32, 3)

test_labels.shape # samples

(26032,)

extra_images.shape # samples, height, width, channels

(531131, 32, 32, 3)

extra_labels.shape # samples


(531131,)

Examine a single image

```
single_image = train_images[25]
single_label = train_labels[25]
```


```
plt.imshow(single_image)
print(f"Label: {single_label}")
```

↻ Label: 3



```
single_image_gray = np.mean(single_image, axis=-1)
plt.imshow(single_image_gray, cmap="gray")
print(f"Label: {single_label}")
```

↻ Label: 3



```
np.unique(train_labels)
```

↻ array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10], dtype=uint8)

```
np.unique(test_labels)
```

↻ array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10], dtype=uint8)

```
np.unique(extra_labels)
```

↻ array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10], dtype=uint8)

```
# Replaces '10' labels with '0' labels
```

```
train_labels[train_labels == 10] = 0
```

```
test_labels[test_labels == 10] = 0
```

```
extra_labels[extra_labels == 10] = 0
```

```
np.unique(train_labels)
```

↻ array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=uint8)

```
np.unique(test_labels)
```

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=uint8)
```

```
np.unique(extra_labels)
```

```
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=uint8)
```

```
# Normalize pixel values
```

```
train_images = train_images / 255.0
```

```
test_images = test_images / 255.0
```

```
extra_images = extra_images / 255.0
```

```
normalized_single_image = train_images[25]
```

```
normalized_single_image
```

```
array([[0.32941176, 0.28627451, 0.30588235],
       [0.31764706, 0.28235294, 0.31372549],
       [0.30980392, 0.2745098 , 0.30980392],
       ...,
       [0.19215686, 0.13333333, 0.15686275],
       [0.19607843, 0.14509804, 0.17647059],
       [0.19215686, 0.12941176, 0.17254902]],

      [[0.3372549 , 0.30980392, 0.32156863],
       [0.33333333, 0.31372549, 0.3372549 ],
       [0.31764706, 0.29803922, 0.3254902 ],
       ...,
       [0.18431373, 0.13333333, 0.14509804],
       [0.18823529, 0.1372549 , 0.16078431],
       [0.18431373, 0.12156863, 0.15686275]],

      [[0.32941176, 0.32156863, 0.32941176],
       [0.3254902 , 0.32156863, 0.34117647],
       [0.27843137, 0.27058824, 0.29803922],
       ...,
       [0.19607843, 0.15294118, 0.15294118],
       [0.19607843, 0.15294118, 0.15686275],
       [0.18823529, 0.13333333, 0.15294118]],

      ...,

      [[0.50980392, 0.53333333, 0.50980392],
       [0.52941176, 0.54509804, 0.53333333],
       [0.52941176, 0.5254902 , 0.5372549 ],
       ...,
       [0.31372549, 0.26666667, 0.29803922],
       [0.31372549, 0.27058824, 0.29411765],
       [0.30196078, 0.26666667, 0.28235294]],


      [[0.43529412, 0.46666667, 0.43921569],
       [0.4745098 , 0.49803922, 0.47058824],
       [0.49019608, 0.49411765, 0.48627451],
       ...,
       [0.3254902 , 0.2745098 , 0.31764706],
       [0.3254902 , 0.2745098 , 0.30980392],
       [0.31372549, 0.27058824, 0.29803922]],

      [[0.34117647, 0.37254902, 0.36078431],
       [0.36862745, 0.39607843, 0.36862745],
       [0.4          , 0.4          , 0.39607843],
       ...,
       [0.3372549 , 0.28235294, 0.32941176],
       [0.33333333, 0.28235294, 0.32156863],
       [0.32156863, 0.2745098 , 0.30588235]]])
```

```
normalized_single_image_gray = np.mean(normalized_single_image, axis=-1)
```

```
plt.imshow(normalized_single_image_gray, cmap="gray")
```

↗ <matplotlib.image.AxesImage at 0x7ac4e9fcadd0>



```
# One-hot encode the labels
cat_train_labels = tf.keras.utils.to_categorical(train_labels, num_classes=10)
cat_test_labels = tf.keras.utils.to_categorical(test_labels, num_classes=10)
cat_extra_labels = tf.keras.utils.to_categorical(extra_labels, num_classes=10)
```

```
train_labels[25]
```

↗ np.uint8(3)

```
cat_train_labels[25]
```

↗ array([0., 0., 0., 1., 0., 0., 0., 0., 0., 0.])

```
# Combine train and extra sets
train_images = np.concatenate((train_images, extra_images), axis=0)
cat_train_labels = np.concatenate((cat_train_labels, cat_extra_labels), axis=0)
```

```
train_images.shape
```

↗ (604388, 32, 32, 3)

```
cat_train_labels.shape
```

↗ (604388, 10)

2) Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.

```
random_indices = np.random.choice(604388, size=20, replace=False)
```

```
random_indices
```

↗ array([497288, 570002, 556502, 183622, 212624, 467273, 47412, 417272,
490139, 141131, 316092, 196512, 4671, 508397, 260117, 16556,
88078, 334103, 237877, 403208])


```
plt.figure(figsize=(15, 4))
```

```
for i, idx in enumerate(random_indices):
    image = train_images[idx, :, :, :] # shape: (32, 32, 3)
    label = cat_train_labels[idx]
```

```
plt.subplot(1, 20, i + 1)
plt.imshow(image)
```

```
plt.title(f'{label}')
plt.axis('off')

plt.tight_layout()
plt.show()
```



3) Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. Hint: retain the channel dimension, which will now have size 1.

```
train_images_gray = np.mean(train_images, axis=-1, keepdims=True) # shape: (123257, 32, 32, 1)
test_images_gray = np.mean(test_images, axis=-1, keepdims=True) # shape: (26032, 32, 32, 1)
```

```
train_images_gray.shape
```

```
(604388, 32, 32, 1)
```

```
test_images_gray.shape
```

```
(26032, 32, 32, 1)
```

4) Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

```
plt.figure(figsize=(15, 4))

for i, idx in enumerate(random_indices):
    image = train_images_gray[idx, :, :, :]
    label = cat_train_labels[idx]

    plt.subplot(1, 20, i + 1)
    plt.imshow(image, cmap="gray")
    plt.title(f'{label}')
    plt.axis('off')

plt.tight_layout()
plt.show()
```



```
from sklearn.model_selection import train_test_split
```

```
# Stratified split (15% validation)
train_imgs, val_imgs, train_lbls, val_lbls = train_test_split(
    train_images_gray,
    cat_train_labels,
    test_size=0.15,
    stratify=np.argmax(cat_train_labels, axis=1),
    random_state=42
)
```

```
train_imgs.shape
```

```
(513729, 32, 32, 1)
```

```
val_imgs.shape
```

```
(90659, 32, 32, 1)
```

```
train_lbls.shape
```

```
(513729, 10)
```

```
val_lbls.shape
```

```
(90659, 10)
```

✓ (40pts) CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output. Please use only what we learned in class.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.*
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track accuracy metric, and use early stopping during training.
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPool2D, Flatten, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping
```

```
model = Sequential()
```

```
# Convolutional Block 1
model.add(Conv2D(filters=32, kernel_size=(3, 3), padding='same', activation='relu', input_shape=(32, 32, 1)))
model.add(MaxPool2D(pool_size=(2, 2)))
```

```
# Convolutional Block 2
model.add(Conv2D(filters=64, kernel_size=(3, 3), padding='same', activation='relu'))
model.add(MaxPool2D(pool_size=(2, 2)))
```

```
# Convolutional Block 3
model.add(Conv2D(filters=128, kernel_size=(3, 3), padding='same', activation='relu'))
model.add(MaxPool2D(pool_size=(2, 2)))
```

```
# Fully Connected Layers
model.add(Flatten())
model.add(Dense(units=256, activation='relu'))
model.add(Dropout(rate=0.5))
model.add(Dense(units=10, activation='softmax'))
```

```
# Compile
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
↳ /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	320
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_1 (Conv2D)	(None, 16, 16, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_2 (Conv2D)	(None, 8, 8, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 256)	524,544
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 10)	2,570

Total params: 619,786 (2.36 MB)
 Trainable params: 619,786 (2.36 MB)
 Non-trainable params: 0 (0.00 B)

Early stop

early_stop = EarlyStopping(monitor='val_loss', patience=3)

Train

model.fit(train_imgs, train_lbls, epochs=30, validation_data=(val_imgs, val_lbls), callbacks=[early_stop])

```

Epoch 1/30
16055/16055 ————— 59s 3ms/step - accuracy: 0.8205 - loss: 0.5608 - val_accuracy: 0.9549 - val_loss: 0.1622
Epoch 2/30
16055/16055 ————— 51s 3ms/step - accuracy: 0.9484 - loss: 0.1826 - val_accuracy: 0.9640 - val_loss: 0.1366
Epoch 3/30
16055/16055 ————— 50s 3ms/step - accuracy: 0.9575 - loss: 0.1535 - val_accuracy: 0.9629 - val_loss: 0.1375
Epoch 4/30
16055/16055 ————— 51s 3ms/step - accuracy: 0.9612 - loss: 0.1412 - val_accuracy: 0.9672 - val_loss: 0.1276
Epoch 5/30
16055/16055 ————— 53s 3ms/step - accuracy: 0.9640 - loss: 0.1326 - val_accuracy: 0.9688 - val_loss: 0.1213
Epoch 6/30
16055/16055 ————— 52s 3ms/step - accuracy: 0.9655 - loss: 0.1278 - val_accuracy: 0.9693 - val_loss: 0.1194
Epoch 7/30
16055/16055 ————— 51s 3ms/step - accuracy: 0.9667 - loss: 0.1240 - val_accuracy: 0.9693 - val_loss: 0.1225
Epoch 8/30
16055/16055 ————— 51s 3ms/step - accuracy: 0.9675 - loss: 0.1200 - val_accuracy: 0.9711 - val_loss: 0.1174
Epoch 9/30
16055/16055 ————— 50s 3ms/step - accuracy: 0.9686 - loss: 0.1153 - val_accuracy: 0.9680 - val_loss: 0.1324
Epoch 10/30
16055/16055 ————— 51s 3ms/step - accuracy: 0.9690 - loss: 0.1162 - val_accuracy: 0.9688 - val_loss: 0.1252
Epoch 11/30
16055/16055 ————— 51s 3ms/step - accuracy: 0.9701 - loss: 0.1137 - val_accuracy: 0.9713 - val_loss: 0.1188
<keras.src.callbacks.history.History at 0x7abdd0228710>

```

model.metrics_names # Check the metric names


```
['loss', 'compile_metrics']
```

losses = pd.DataFrame(model.history.history)


losses

	accuracy	loss	val_accuracy	val_loss
0	0.902721	0.324831	0.954930	0.162200
1	0.950602	0.176437	0.963997	0.136565
2	0.957807	0.153704	0.962861	0.137525
3	0.960828	0.142849	0.967196	0.127583
4	0.963562	0.135306	0.968762	0.121255
5	0.965013	0.129435	0.969347	0.119449
6	0.966064	0.125630	0.969292	0.122537
7	0.967095	0.121628	0.971078	0.117439
8	0.967859	0.119363	0.968045	0.132404
9	0.968442	0.118479	0.968795	0.125223
10	0.968877	0.117038	0.971266	0.118845

```
losses[['accuracy', 'val_accuracy']].plot()
```



```
losses[['loss', 'val_loss']].plot()
```



```
from sklearn.metrics import classification_report, confusion_matrix
```

```
predictions = model.predict(test_images_gray)
```

↗ 814/814 ————— 2s 2ms/step

```
predictions
```

```
↗ array([[4.8619563e-11, 8.7852543e-09, 2.4123054e-07, ..., 1.6935695e-07,
          3.4333281e-09, 1.4450826e-08],
         [1.2798432e-09, 7.1566602e-09, 9.9999416e-01, ..., 2.2324438e-07,
          9.0855612e-10, 3.6524348e-09],
         [1.4622789e-08, 9.9998963e-01, 6.7177808e-07, ..., 9.8884368e-07,
          1.3059671e-08, 1.5191265e-08],
         ...,
         [6.8170480e-10, 1.3964177e-07, 4.4802335e-05, ..., 9.9995232e-01,
          3.3615491e-12, 3.8908124e-07],
         [1.2936059e-05, 7.7630782e-09, 1.3935372e-10, ..., 2.3640118e-10,
          1.5215099e-05, 1.1328984e-11],
         [1.9599569e-12, 1.4076382e-05, 4.7574801e-07, ..., 9.9998546e-01,
          5.5827988e-13, 1.5650263e-09]], dtype=float32)
```

```
pred_binary = predictions > 0.5
```

```
pred_binary
```

```
↗ array([[False, False, False, ..., False, False, False],
         [False, False, True, ..., False, False, False],
         [False, True, False, ..., False, False, False],
         ...,
         [False, False, False, ..., True, False, False],
         [False, False, False, ..., False, False, False],
         [False, False, False, ..., True, False, False]])
```

```
pred_class = np.argmax(predictions, axis=1)
```

```
pred_class
```

```
↗ array([5, 2, 1, ..., 7, 6, 7])
```

```
cat_test_labels[0]
```

```
↗ array([0., 0., 0., 0., 0., 1., 0., 0., 0., 0.])
```

```
print(classification_report(test_labels,pred_class))
```

```
↗
```


	precision	recall	f1-score	support
0	0.92	0.96	0.94	1744
1	0.95	0.97	0.96	5099
2	0.97	0.95	0.96	4149
3	0.92	0.94	0.93	2882
4	0.97	0.95	0.96	2523
5	0.95	0.94	0.94	2384
6	0.93	0.94	0.94	1977
7	0.96	0.94	0.95	2019
8	0.94	0.91	0.93	1660
9	0.93	0.93	0.93	1595
accuracy			0.95	26032
macro avg	0.94	0.94	0.94	26032
weighted avg	0.95	0.95	0.95	26032

```
confusion_matrix(test_labels,pred_class)
```

```
↗ array([[1669, 13, 6, 7, 2, 4, 16, 5, 5, 17],
         [ 45, 4930, 22, 14, 21, 10, 12, 29, 10, 6],
         [ 8, 37, 3961, 64, 14, 17, 5, 21, 10, 12],
         [ 7, 40, 17, 2721, 4, 35, 13, 4, 19, 22],
         [ 15, 59, 8, 16, 2386, 5, 14, 8, 3, 9],
         [ 5, 10, 12, 59, 4, 2246, 26, 2, 11, 9],
         [ 28, 17, 3, 12, 5, 24, 1855, 3, 24, 6],
         [ 3, 69, 18, 11, 5, 9, 5, 1890, 1, 8],
         [ 17, 7, 6, 28, 8, 11, 41, 2, 1513, 27],
         [ 27, 12, 22, 14, 8, 10, 1, 2, 8, 1491]])
```

```
import seaborn as sns
```


```
sns.heatmap(confusion_matrix(test_labels,pred_class),annot=True)
```



```
my_number = test_images_gray[664]
```

```
plt.imshow(my_number.reshape(32, 32), cmap='gray')
```

<matplotlib.image.AxesImage at 0x7abc97558090>



```
model.predict(my_number.reshape(1, 32, 32, 1))
```

1/1 ————— 0s 401ms/step
 array([[6.4053495e-07, 2.6309046e-05, 3.0764255e-05, 2.8490178e-06,
 9.9993253e-01, 6.6887559e-08, 1.5111578e-06, 4.9039277e-06,
 5.2862475e-08, 3.4097283e-07]], dtype=float32)

```
np.round(model.predict(my_number.reshape(1, 32, 32, 1)), 2)
```

1/1 ————— 0s 30ms/step
 array([[0., 0., 0., 0., 1., 0., 0., 0., 0., 0.]], dtype=float32)

```
model.save('base_model_full_dataset.keras')
```

✓ (30pts) Improvement

- Try to use ImageDataGenerator to augment the data.
- Try to add BatchNormalization layers using Keras API. Please do a little research how this is done.

- Try to use a pre-trained CNN model for transfer learning. You make your choice and explain why.

Does any of the above help?

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```


```
#help(ImageDataGenerator)
```

```
train_gen = ImageDataGenerator(
    rotation_range=10,
    width_shift_range=0.1,
    height_shift_range=0.1,
    zoom_range=0.1,
    shear_range=8,
    fill_mode='nearest'
)
```

```
pass_gen = ImageDataGenerator()
```


```
plt.imshow(train_images[0])
```

```
<matplotlib.image.AxesImage at 0x7abc9749e0d0>
```



```
plt.imshow(train_gen.random_transform(train_images[0]))
```

```
<matplotlib.image.AxesImage at 0x7abb6433d6d0>
```



BATCH = 64

```
train_gen = train_gen.flow(
    train_imgs, # numpy array (N,32,32,1)
    train_lbls, # numpy array (N,10)
    batch_size=BATCH,
    shuffle=True
)
```

```
val_gen = pass_gen.flow(
    val_imgs,
    val_lbls,
    batch_size=BATCH,
    shuffle=False
)
```

```
test_gen = pass_gen.flow(
    test_images_gray,
    cat_test_labels,
    batch_size=BATCH,
    shuffle=False
)
```

```
print("Samples :", train_gen.n)
print("Batch sz:", train_gen.batch_size)
```

```
batch_x, batch_y = train_gen[0]
print("Batch X shape:", batch_x.shape)
print("Batch Y shape:", batch_y.shape)
```

```
↗ Samples : 513729
  Batch sz: 64
  Batch X shape: (64, 32, 32, 1)
  Batch Y shape: (64, 10)
```

```
model2 = Sequential()
```

```
# Convolutional Block 1
model2.add(Conv2D(filters=32, kernel_size=(3, 3), padding='same', activation='relu', input_shape=(32, 32, 1)))
model2.add(MaxPool2D(pool_size=(2, 2)))
```

```
# Convolutional Block 2
model2.add(Conv2D(filters=64, kernel_size=(3, 3), padding='same', activation='relu'))
model2.add(MaxPool2D(pool_size=(2, 2)))
```

```
# Convolutional Block 3
model2.add(Conv2D(filters=128, kernel_size=(3, 3), padding='same', activation='relu'))
model2.add(MaxPool2D(pool_size=(2, 2)))
```

```
# Fully Connected Layers
model2.add(Flatten())
model2.add(Dense(units=256, activation='relu'))
model2.add(Dropout(rate=0.5))
model2.add(Dense(units=10, activation='softmax'))
```

```
# Compile
model2.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
↗ /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
# Train
model2.fit(
    train_gen,
    epochs=30,
    validation_data=val_gen,
    callbacks=[early_stop]
)
```

```
↗ Epoch 1/30
  /usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDat
    self._warn_if_super_not_called()
  8028/8028 ————— 154s 19ms/step - accuracy: 0.7230 - loss: 0.8355 - val_accuracy: 0.9559 - val_loss: 0.1594
  Epoch 2/30
  8028/8028 ————— 150s 19ms/step - accuracy: 0.9193 - loss: 0.2731 - val_accuracy: 0.9666 - val_loss: 0.1258
  Epoch 3/30
  8028/8028 ————— 151s 19ms/step - accuracy: 0.9344 - loss: 0.2275 - val_accuracy: 0.9675 - val_loss: 0.1251
```

```

Epoch 4/30
8028/8028 ————— 152s 19ms/step - accuracy: 0.9415 - loss: 0.2026 - val_accuracy: 0.9671 - val_loss: 0.1223
Epoch 5/30
8028/8028 ————— 155s 19ms/step - accuracy: 0.9470 - loss: 0.1869 - val_accuracy: 0.9693 - val_loss: 0.1182
Epoch 6/30
8028/8028 ————— 153s 19ms/step - accuracy: 0.9492 - loss: 0.1784 - val_accuracy: 0.9717 - val_loss: 0.1104
Epoch 7/30
8028/8028 ————— 153s 19ms/step - accuracy: 0.9506 - loss: 0.1758 - val_accuracy: 0.9726 - val_loss: 0.1092
Epoch 8/30
8028/8028 ————— 159s 20ms/step - accuracy: 0.9532 - loss: 0.1668 - val_accuracy: 0.9726 - val_loss: 0.1065
Epoch 9/30
8028/8028 ————— 153s 19ms/step - accuracy: 0.9538 - loss: 0.1653 - val_accuracy: 0.9727 - val_loss: 0.1074
Epoch 10/30
8028/8028 ————— 152s 19ms/step - accuracy: 0.9552 - loss: 0.1598 - val_accuracy: 0.9737 - val_loss: 0.1043
Epoch 11/30
8028/8028 ————— 152s 19ms/step - accuracy: 0.9564 - loss: 0.1586 - val_accuracy: 0.9737 - val_loss: 0.1055
Epoch 12/30
8028/8028 ————— 153s 19ms/step - accuracy: 0.9572 - loss: 0.1563 - val_accuracy: 0.9748 - val_loss: 0.1034
Epoch 13/30
8028/8028 ————— 152s 19ms/step - accuracy: 0.9571 - loss: 0.1530 - val_accuracy: 0.9730 - val_loss: 0.1035
Epoch 14/30
8028/8028 ————— 153s 19ms/step - accuracy: 0.9590 - loss: 0.1492 - val_accuracy: 0.9751 - val_loss: 0.0984
Epoch 15/30
8028/8028 ————— 149s 19ms/step - accuracy: 0.9589 - loss: 0.1494 - val_accuracy: 0.9755 - val_loss: 0.0993
Epoch 16/30
8028/8028 ————— 148s 18ms/step - accuracy: 0.9590 - loss: 0.1482 - val_accuracy: 0.9736 - val_loss: 0.1005
Epoch 17/30
8028/8028 ————— 147s 18ms/step - accuracy: 0.9596 - loss: 0.1464 - val_accuracy: 0.9739 - val_loss: 0.1039
<keras.src.callbacks.history.History at 0x7abb64328710>

```

```
model2.metrics_names # Check the metric names
```

```
['loss', 'compile_metrics']
```

```
losses2 = pd.DataFrame(model2.history.history)
```


```
losses2
```

```


accuracy    loss  val_accuracy  val_loss
0    0.842296  0.502098      0.955912  0.159390
1    0.923759  0.260971      0.966578  0.125800
2    0.936525  0.219445      0.967494  0.125083
3    0.942843  0.200806      0.967108  0.122278
4    0.946651  0.188646      0.969303  0.118165
5    0.949501  0.179973      0.971729  0.110425
6    0.950959  0.173497      0.972579  0.109186
7    0.953220  0.167676      0.972623  0.106470
8    0.953822  0.163964      0.972678  0.107391
9    0.955315  0.160435      0.973660  0.104333
10   0.955870  0.159602      0.973682  0.105471
11   0.956759  0.156551      0.974796  0.103416
12   0.957026  0.154638      0.972987  0.103511
13   0.958420  0.152850      0.975116  0.098370
14   0.958394  0.151383      0.975469  0.099337
15   0.958953  0.149758      0.973604  0.100506
16   0.959533  0.148404      0.973924  0.103896

```

```
losses2[['accuracy', 'val_accuracy']].plot()
```



```
losses2[['loss', 'val_loss']].plot()
```



```
predictions2 = model2.predict(test_images_gray)
```

814/814 ————— 2s 2ms/step

```
predictions2
```

```
array([[5.2060019e-09, 4.3295935e-07, 1.9128575e-07, ..., 5.7688095e-07,
        2.2260602e-09, 2.5183297e-07],
       [9.9598385e-11, 4.6322217e-08, 9.9999964e-01, ..., 1.1171285e-07,
        1.4305229e-10, 2.2930874e-10],
       [2.7705939e-06, 9.9817336e-01, 1.3308752e-06, ..., 1.1998583e-05,
        2.5519222e-07, 4.2033093e-07],
       ...,
       [6.9868920e-09, 2.2034372e-05, 7.8806333e-06, ..., 9.9996006e-01,
        1.6232604e-10, 1.9304376e-07],
       [5.6024688e-05, 4.0376702e-07, 2.0019315e-09, ..., 5.9883138e-09,
        1.3196832e-06, 3.2170625e-08],
       [6.8687855e-10, 3.7536487e-05, 7.8305834e-07, ..., 9.9996173e-01,
        1.6771548e-12, 6.2388050e-09]], dtype=float32)
```

```
pred2_binary = predictions2 > 0.5
```

```
pred2_binary
```

```
array([[False, False, False, ..., False, False, False],
       [False, False, True, ..., False, False, False],
```

```
[False, True, False, ..., False, False, False],
...,
[False, False, False, ..., True, False, False],
[False, False, False, ..., False, False, False],
[False, False, False, ..., True, False, False]])
```

```
pred2_class = np.argmax(predictions2, axis=1)
```

```
pred2_class
```

```
array([5, 2, 1, ..., 7, 6, 7])
```

```
print(classification_report(test_labels, pred2_class))
```

```

      precision    recall  f1-score   support

     0       0.93      0.96      0.95       1744
     1       0.97      0.96      0.97       5099
     2       0.97      0.98      0.97       4149
     3       0.94      0.95      0.94       2882
     4       0.96      0.97      0.97       2523
     5       0.96      0.95      0.95       2384
     6       0.95      0.95      0.95       1977
     7       0.96      0.95      0.96       2019
     8       0.95      0.92      0.94       1660
     9       0.95      0.92      0.94       1595


 accuracy          0.95
 macro avg          0.95
 weighted avg       0.96
```

```
confusion_matrix(test_labels, pred2_class)
```

```

array([[1679, 10, 7, 8, 1, 7, 21, 4, 0, 7],
       [ 22, 4915, 25, 27, 38, 11, 11, 37, 7, 6],
       [ 5, 16, 4056, 16, 10, 8, 5, 24, 4, 5],
       [ 5, 20, 24, 2729, 6, 33, 10, 6, 21, 28],
       [ 6, 28, 16, 5, 2447, 3, 3, 4, 3, 8],
       [ 5, 14, 6, 54, 5, 2268, 19, 1, 9, 3],
       [ 21, 13, 4, 9, 5, 23, 1875, 3, 23, 1],
       [ 2, 55, 20, 11, 6, 2, 0, 1919, 1, 3],
       [ 15, 5, 13, 25, 11, 6, 36, 0, 1529, 20],
       [ 38, 10, 23, 17, 9, 9, 3, 1, 11, 1474]])
```


```
sns.heatmap(confusion_matrix(test_labels, pred2_class), annot=True)
```



```
my_number = test_images_gray[1000]
```

```
plt.imshow(my_number.reshape(32, 32), cmap='gray')
```


↗ <matplotlib.image.AxesImage at 0x7ac14dd7aa90>



```
model2.predict(my_number.reshape(1, 32, 32, 1))
```

↗ 1/1 ————— 0s 247ms/step
array([[4.51854554e-10, 2.19837915e-07, 8.16593959e-10, 9.18656099e-07,
1.22803447e-08, 9.99963164e-01, 3.54813783e-05, 7.46057272e-10,
8.70424888e-09, 1.09917956e-07]], dtype=float32)

```
np.round(model2.predict(my_number.reshape(1, 32, 32, 1)), 2)
```

↗ 1/1 ————— 0s 33ms/step
array([[0., 0., 0., 0., 0., 1., 0., 0., 0., 0.]], dtype=float32)

```
model2.save('data_augmentation_model_full_dataset.keras')
```

Adding Batch Normalization

```
from tensorflow.keras.layers import BatchNormalization, Activation
```

```
model3 = Sequential()
```

```
# Convolutional Block 1
```

```
model3.add(Conv2D(filters=32, kernel_size=(3, 3), padding='same', use_bias=False, input_shape=(32, 32, 1)))  
model3.add(BatchNormalization())  
model3.add(Activation('relu'))  
model3.add(MaxPool2D(pool_size=(2, 2)))
```

```
# Convolutional Block 2
```

```
model3.add(Conv2D(filters=64, kernel_size=(3, 3), padding='same', use_bias=False))  
model3.add(BatchNormalization())  
model3.add(Activation('relu'))  
model3.add(MaxPool2D((2, 2)))
```

```
# Convolutional Block 3
```

```
model3.add(Conv2D(filters=128, kernel_size=(3, 3), padding='same', use_bias=False))  
model3.add(BatchNormalization())  
model3.add(Activation('relu'))  
model3.add(MaxPool2D((2, 2)))
```

```
# Fully Connected Layers
```

```
model3.add(Flatten())  
model3.add(Dense(units=256, activation='relu'))  
model3.add(Dropout(rate=0.5))  
model3.add(Dense(units=10, activation='softmax'))
```

```
# Compile
```

```
model3.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

↗ /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```
# Train
model3.fit(
    train_gen,
    epochs=30,
    validation_data=val_gen,
    callbacks=[early_stop]
)
```

```
Epoch 1/30
8028/8028 ————— 162s 20ms/step - accuracy: 0.4983 - loss: 1.3425 - val_accuracy: 0.9453 - val_loss: 0.2114
Epoch 2/30
8028/8028 ————— 154s 19ms/step - accuracy: 0.8592 - loss: 0.4366 - val_accuracy: 0.9572 - val_loss: 0.1576
Epoch 3/30
8028/8028 ————— 162s 20ms/step - accuracy: 0.9382 - loss: 0.2218 - val_accuracy: 0.9713 - val_loss: 0.1136
Epoch 4/30
8028/8028 ————— 154s 19ms/step - accuracy: 0.9516 - loss: 0.1760 - val_accuracy: 0.9691 - val_loss: 0.1176
Epoch 5/30
8028/8028 ————— 153s 19ms/step - accuracy: 0.9569 - loss: 0.1581 - val_accuracy: 0.9733 - val_loss: 0.1059
Epoch 6/30
8028/8028 ————— 152s 19ms/step - accuracy: 0.9604 - loss: 0.1464 - val_accuracy: 0.9722 - val_loss: 0.1056
Epoch 7/30
8028/8028 ————— 153s 19ms/step - accuracy: 0.9618 - loss: 0.1398 - val_accuracy: 0.9744 - val_loss: 0.1000
Epoch 8/30
8028/8028 ————— 155s 19ms/step - accuracy: 0.9642 - loss: 0.1336 - val_accuracy: 0.9731 - val_loss: 0.1033
Epoch 9/30
8028/8028 ————— 153s 19ms/step - accuracy: 0.9654 - loss: 0.1288 - val_accuracy: 0.9638 - val_loss: 0.1329
Epoch 10/30
8028/8028 ————— 153s 19ms/step - accuracy: 0.9667 - loss: 0.1237 - val_accuracy: 0.9741 - val_loss: 0.0987
Epoch 11/30
8028/8028 ————— 151s 19ms/step - accuracy: 0.9670 - loss: 0.1239 - val_accuracy: 0.9754 - val_loss: 0.0949
Epoch 12/30
8028/8028 ————— 150s 19ms/step - accuracy: 0.9679 - loss: 0.1210 - val_accuracy: 0.9737 - val_loss: 0.1003
Epoch 13/30
8028/8028 ————— 148s 18ms/step - accuracy: 0.9685 - loss: 0.1192 - val_accuracy: 0.9749 - val_loss: 0.0965
Epoch 14/30
8028/8028 ————— 149s 19ms/step - accuracy: 0.9690 - loss: 0.1164 - val_accuracy: 0.9714 - val_loss: 0.1066
<keras.src.callbacks.history.History at 0x7ac4ea04a990>
```

```
model3.metrics_names # Check the metric names
```

```
['loss', 'compile_metrics']
```


```
losses3 = pd.DataFrame(model3.history.history)
```

```
losses3
```

```
accuracy    loss  val_accuracy  val_loss
0    0.625666  1.002552    0.945334  0.211423
1    0.891234  0.353513    0.957158  0.157647
2    0.942843  0.207207    0.971332  0.113594
3    0.952714  0.172306    0.969126  0.117628
4    0.957357  0.156275    0.973284  0.105890
5    0.960697  0.144803    0.972181  0.105557
6    0.962375  0.138645    0.974365  0.100029
7    0.964549  0.132129    0.973141  0.103281
8    0.965544  0.128588    0.963798  0.132867
9    0.966360  0.125442    0.974123  0.098695
10   0.967125  0.122985    0.975435  0.094902
11   0.967582  0.120684    0.973671  0.100347
12   0.968552  0.117752    0.974939  0.096470
13   0.968941  0.116763    0.971420  0.106648
```


```
losses3[['accuracy', 'val_accuracy']].plot()
```

<Axes: >



losses3[['loss', 'val_loss']].plot()

<Axes: >



predictions3 = model3.predict(test_images_gray)

814/814 ————— 2s 2ms/step

predictions3

```
array([[2.8453124e-07, 2.8197126e-05, 8.7615354e-06, ..., 1.5913520e-05,
        5.3405046e-07, 4.4351059e-06],
       [8.0707914e-06, 6.1902298e-05, 9.9964201e-01, ..., 4.7926307e-05,
        2.4544610e-05, 3.5739791e-05],
       [6.1016585e-06, 9.9915576e-01, 2.3376668e-05, ..., 6.3453408e-05,
        1.7735758e-05, 1.9348902e-06],
       ...,
       [5.8625133e-08, 8.8640722e-05, 1.0469111e-04, ..., 9.9977797e-01,
        3.2194031e-08, 1.8442601e-05],
       [1.5912355e-03, 4.6303289e-04, 4.1217943e-05, ..., 5.1023730e-05,
        1.8465503e-04, 3.5647965e-05],
       [7.2829508e-08, 1.1128939e-03, 1.3698303e-05, ..., 9.9885786e-01,
        2.6150971e-08, 8.8013776e-06]], dtype=float32)
```

pred3_binary = predictions3 > 0.5

pred3_binary

```
array([[False, False, False, ..., False, False, False],
       [False, False, True, ..., False, False, False],
```

```
[False, True, False, ..., False, False, False],
...,
[False, False, False, ..., True, False, False],
[False, False, False, ..., False, False, False],
[False, False, False, ..., True, False, False]])
```

```
pred3_class = np.argmax(predictions3, axis=1)
```

```
pred3_class
```

```
array([5, 2, 1, ..., 7, 6, 7])
```

```
print(classification_report(test_labels, pred3_class))
```

```
precision    recall  f1-score   support


0           0.95      0.94      0.95     1744
1           0.95      0.97      0.96     5099
2           0.97      0.96      0.97     4149
3           0.93      0.94      0.94     2882
4           0.96      0.96      0.96     2523
5           0.95      0.95      0.95     2384
6           0.94      0.94      0.94     1977
7           0.95      0.93      0.94     2019
8           0.96      0.91      0.94     1660
9           0.94      0.95      0.94     1595

accuracy          0.95
macro avg         0.95
weighted avg      0.95
```

```
confusion_matrix(test_labels, pred3_class)
```


```
array([[1645, 16, 11, 13, 3, 7, 30, 2, 0, 17],
       [18, 4938, 17, 36, 41, 10, 6, 27, 3, 3],
       [4, 39, 3999, 22, 17, 4, 3, 47, 5, 9],
       [7, 27, 19, 2718, 10, 39, 8, 8, 17, 29],
       [6, 44, 12, 11, 2429, 4, 5, 7, 2, 3],
       [3, 8, 7, 59, 7, 2265, 19, 3, 5, 8],
       [17, 12, 9, 9, 9, 25, 1866, 2, 24, 4],
       [4, 99, 14, 4, 3, 3, 1, 1885, 1, 5],
       [10, 12, 11, 24, 9, 9, 39, 4, 1518, 24],
       [14, 10, 20, 13, 7, 10, 1, 1, 6, 1513]])
```

```
sns.heatmap(confusion_matrix(test_labels, pred3_class), annot=True)
```



```
my_number = test_images_gray[3333]
```

```
plt.imshow(my_number.reshape(32, 32), cmap='gray')
```

 <matplotlib.image.AxesImage at 0x7abac4748f50>

